Jacksonville University — School of Engineering & Technology

CS302 — Software Design & Development: AI Engineering — Fall 2026

M2 Governance Appendix — Template

Team name:	**Date:**
Team Tokenizer (CuraAI)	October 21, 2026

Attach this appendix to your M2 design document. Complete all five governance ADRs — one row set each. Keep each field to a sentence or two; the full context/options/consequences reasoning belongs in your numbered ADR trail (reference the ADR number here). See the CS302 AIGRF handout for what reviewers look for.

Governance ADR 1 — Decision Transparency

How does every AI decision in your system expose its reasoning path — inputs, context, output, and who or what accepted it?

Decision	Every AI call is logged: what went in, what came out, and who accepted it. Doctors see the source for every item. (Governance ADR 1)
Wired where in Main	Main wraps every model call with a logger that writes to the audit log.
Evidence / metric	Every test call has a matching log entry; every review row shows its source.
Owner (steward)	Nicole Hood

Governance ADR 2 — Observability

Which run logs and metrics (off-ramp rate, failed-guardrail count, …) are designed in, and where does the logging wrapper live?

Decision	We log every case and count handoffs to the doctor, failed AI checks, follow-up questions, time, and cost. (Governance ADR 2)
Wired where in Main	The same logger wraps the model calls; Main creates it once and shares it.
Evidence / metric	Numbers reported per 100 test cases; alert limits are set from that baseline.
Owner (steward)	RJ Peck

Governance ADR 3 — HITL (Human-In-The-Loop) Gates

Which paths require human approval (money / safety / employment / external comms), and who approves?

Decision	The doctor approves every clinical decision, prescription, and patient reply. AI only asks pre-approved follow-up questions and fetches records. (Governance ADR 3)
Wired where in Main	Main gives the AI no write access to the EHR, so it cannot change the chart or send free text.
Evidence / metric	A test confirms no write function exists; the log shows no AI orders or free-text messages.
Owner (steward)	Bryan Ponce

Governance ADR 4 — Tool Least-Privilege

What can each agent tool reach, per its allow-list and scope — and what is deliberately unreachable?

Decision	The AI sees only symptom text. EHR and drug access is read-only, one patient, 3 weeks back. It cannot write, order, or see other patients. (Governance ADR 4)

Wired where in Main	Main connects only read-only adapters and a short allow-list of lookups.
Evidence / metric	Tests show out-of-scope requests are rejected, including prompt-injection attempts.
Owner (steward)	Downing Nightingale

Governance ADR 5 — Cost Control

What are the per-case caps and query budgets, and what trips the circuit-breaker?

Decision	Each case has limits (proposed: 6 AI calls, 5 lookups, 2 follow-ups, $0.25). If a limit is hit, the AI stops and the fixed chart pull still runs. (Governance ADR 5)
Wired where in Main	Main wraps the AI calls in a cost guard that checks the limits before every call.
Evidence / metric	Cost per case compared with the $0.10 estimate; tests force each limit to trip.
Owner (steward)	Nicole Hood

Steward sign-off

We can defend each decision above out loud, without notes, as required at the M4 interrogation.

Steward of record (signature):	**Reviewed by (teammate signature):**

AI produces outputs — humans own outcomes.